

Exploratory data analysis

Coffee bean quality dataset

Contents

- Introduction
- Aim of the project
- Project workflow
- Data understanding
- Data Cleaning
- Data Filtering
- Univariate Analysis
- Bivariate Analysis
- Multivariate Analysis
- Insight obtained from the Analysis
- Conclusion

Introduction

The dataset used in this project is based on the Coffee Quality Institute (CQI) Coffee Quality Database and contains 1339 Arabica coffee bean reviews performed by the Coffee Quality Institute, it seeks to predict the overall quality score based on intrinsic features of the coffee beans (such

as color, variety, defects), geographical features and information related to the producer, exporter, or grower of the beans.

Aim of the Project

- The primary objective of this project is to conduct a comprehensive analysis to understand the factors that contribute to Coffee quality and improve overall Quality of Coffee .
- By systematically evaluating various parameters such as flavor, aroma, Body, consistency etc., the project aims to identify key factors influencing coffee quality.
- This analysis will provide insights into optimizing Variety of Coffees, Processing Methods etc., and meeting consumer expectations more effectively.

Project Workflow:

Data Understanding

Data preparation is the process of preparing raw data so that it is suitable for further processing and analysis. Key steps include collecting, cleaning, and labeling raw data into a form suitable for further process and then exploring and visualizing the data.

The given coffe quality dataset has been read in the jupyter notebook

```
[302]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
[304]: data=pd.read_csv('C:/Users/Rehana/Desktop/coffeeQuality.csv')
```

```
[305]: data
```

	Unnamed: 0	Species	Owner	Country.of.Origin	Farm.Name	Lot.Number	Mill	ICO.Number	Company	Altitude	...	Color	Category.Two.Defects	E
0	0	Arabica	metad plc	Ethiopia	metad plc	NaN	metad plc	2014/2015	metad agricultural developmet plc	1950-2200	...	Green	0	
1	1	Arabica	metad plc	Ethiopia	metad plc	NaN	metad plc	2014/2015	metad agricultural developmet plc	1950-2200	...	Green	1	
2	2	Arabica	grounds for health admin	Guatemala	san marcos barrancas "san cristobal cuch	NaN	NaN	NaN	NaN	1600 - 1800 m	...	NaN	0	
3	3	Arabica	yidnekachew dabessa	Ethiopia	yidnekachew dabessa coffee plantation	NaN	wolensu	NaN	yidnekachew dabessa coffee plantation	1800-2200	...	Green	2	
4	4	Arabica	metad plc	Ethiopia	metad plc	NaN	metad plc	2014/2015	metad agricultural developmet plc	1950-2200	...	Green	2	
...

```
[306]: data.shape
```

```
[306]: (1339, 44)
```

The above code is used to define the rows and columns in the given dataset.

```
[308]: data.tail()
```

```
[308]:
```

	Unnamed: 0	Species	Owner	Country.of.Origin	Farm.Name	Lot.Number	Mill	ICO.Number	Company	Altitude	...	Color	Category.Two.Defects	Expirati
1334	1334	Robusta	luis robles	Ecuador	robustasa	Lavado 1	our own lab	NaN	robustasa	NaN	...	Blue-Green	1	Janua 18th, 20
1335	1335	Robusta	luis robles	Ecuador	robustasa	Lavado 3	own laboratory	NaN	robustasa	40	...	Blue-Green	0	Janua 18th, 20
1336	1336	Robusta	james moore	United States	fazenda cazengo	NaN	cafe cazengo	NaN	global opportunity fund	795 meters	...	NaN	6	Decemb 23rd, 20
1337	1337	Robusta	cafe politico	India	NaN	NaN	NaN	14-1118-2014-0087	cafe politico	NaN	...	Green	1	Augu 25th, 20
1338	1338	Robusta	cafe politico	Vietnam	NaN	NaN	NaN	NaN	cafe politico	NaN	...	NaN	9	Augu 25th, 20

5 rows x 44 columns

```
data.head()
```

	Unnamed: 0	Species	Owner	Country.of.Origin	Farm.Name	Lot.Number	Mill	ICO.Number	Company	Altitude	...	Color	Category.Two.Defects	Expirati
0	0	Arabica	metad plc	Ethiopia	metad plc	NaN	metad plc	2014/2015	metad agricultural developmet plc	1950-2200	...	Green	0	April 2
1	1	Arabica	metad plc	Ethiopia	metad plc	NaN	metad plc	2014/2015	metad agricultural developmet plc	1950-2200	...	Green	1	April 2
2	2	Arabica	grounds for health admin	Guatemala	san marcos barrancas "san cristobal cuch	NaN	NaN	NaN	NaN	1600 - 1800 m	...	NaN	0	May 3 2
3	3	Arabica	yidnekachew dabessa	Ethiopia	yidnekachew dabessa coffee plantation	NaN	wolensu	NaN	yidnekachew dabessa coffee plantation	1800-2200	...	Green	2	Ma 25th, 2
4	4	Arabica	metad plc	Ethiopia	metad plc	NaN	metad plc	2014/2015	metad agricultural developmet plc	1950-2200	...	Green	2	April 2

5 rows x 44 columns

The `info()` method prints information about the DataFrame. The information contains the number of columns, column labels, column data types, memory usage, range index, and the number of cells in each column (non-null values).

```
[309]: data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1339 entries, 0 to 1338
Data columns (total 44 columns):
#   Column              Non-Null Count  Dtype
---  -
0   Unnamed: 0          1339 non-null   int64
1   Species             1339 non-null   object
2   Owner               1332 non-null   object
3   Country.of.Origin   1338 non-null   object
4   Farm.Name           980 non-null    object
5   Lot.Number          276 non-null    object
6   Mill                1021 non-null   object
7   ICO.Number          1180 non-null   object
8   Company             1130 non-null   object
9   Altitude            1113 non-null   object
10  Region              1280 non-null   object
11  Producer            1107 non-null   object
12  Number.of.Bags       1338 non-null   float64
13  Bag.Weight           1339 non-null   object
14  In.Country.Partner   1339 non-null   object
15  Harvest.Year         1292 non-null   object
16  Grading.Date         1339 non-null   object
17  Owner.1              1332 non-null   object
18  Variety             1113 non-null   object
19  Processing.Method    1169 non-null   object
20  Aroma               1339 non-null   float64
21  Flavor              1339 non-null   float64
22  Aftertaste          1339 non-null   float64
23  Acidity              1339 non-null   float64
24  Body                1339 non-null   float64
25  Balance              1339 non-null   float64
26  Uniformity           1339 non-null   float64
27  Clean.Cup            1339 non-null   float64
28  Sweetness            1339 non-null   float64
29  Cupper.Points        1339 non-null   float64
30  Total.Cup.Points     1339 non-null   float64
31  Moisture             1339 non-null   float64
32  Category.One.Defects 1339 non-null   int64
33  Quakers              1338 non-null   float64
34  Color                1069 non-null   object
35  Category.Two.Defects 1339 non-null   int64
36  Expiration           1339 non-null   object
37  Certification.Body    1339 non-null   object
38  Certification.Address 1339 non-null   object
39  Certification.Contact 1339 non-null   object
40  unit_of_measurement  1339 non-null   object
41  altitude_low_meters  1109 non-null   float64
42  altitude_high_meters 1109 non-null   float64
43  altitude_mean_meters 1109 non-null   float64
dtypes: float64(17), int64(3), object(24)
memory usage: 460.4+ KB
```

`Data.describe()` is used to view some basic statistical details like percentile, mean, std, etc. of a data frame or a series of numeric values. The output is shown in the examples below.

```
: data.describe()
```

	Unnamed: 0	Number.of.Bags	Aroma	Flavor	Aftertaste	Acidity	Body	Balance	Uniformity	Clean.Cup	Sweetness	Cupper.Points
count	1339.000000	1338.000000	1339.000000	1339.000000	1339.000000	1339.000000	1339.000000	1339.000000	1339.000000	1339.000000	1339.000000	1339.000000
mean	669.000000	159.085202	7.770187	7.520426	7.401083	7.535706	7.517498	7.518013	9.834877	9.835108	9.856692	7.50337
std	386.680316	173.698167	5.534440	0.398442	0.404463	0.379827	0.370064	0.408943	0.554591	0.763946	0.616102	0.47346
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	334.500000	14.000000	7.420000	7.330000	7.250000	7.330000	7.330000	7.330000	10.000000	10.000000	10.000000	7.25000
50%	669.000000	175.000000	7.580000	7.580000	7.420000	7.580000	7.500000	7.500000	10.000000	10.000000	10.000000	7.50000
75%	1003.500000	275.000000	7.750000	7.750000	7.580000	7.750000	7.670000	7.750000	10.000000	10.000000	10.000000	7.75000
max	1338.000000	3200.000000	200.000000	8.830000	8.670000	8.750000	8.580000	8.750000	10.000000	10.000000	10.000000	10.00000

Data Cleaning

Data cleaning involves revising, rectifying, and organizing information in a dataset to make it consistent and ready for analysis. This step entails identifying and addressing errors, inconsistencies, duplicates, or incomplete entries within the data. The main objective of data cleaning is to enhance the data's quality and usefulness, thereby leading to more dependable and precise findings.


Data. isnull(). sum() returns the number of missing values in the dataset.

```
[312]: data.isnull().sum()
```


```
[312]: Unnamed: 0      0
      Species      0
      Owner        7
      Country.of.Origin  1
      Farm.Name    359
      Lot.Number  1063
      Mill        318
      ICO.Number  159
      Company     209
      Altitude   226
      Region      59
      Producer    232
      Number.of.Bags  1
      Bag.Weight   0
      In.Country.Partner  0
      Harvest.Year  47
      Grading.Date  0
      Owner.1      7
      Variety     226
      Processing.Method 170
      Aroma        0
      Flavor       0
      Aftertaste   0
      Acidity      0
      Body         0
      Balance      0
      Uniformity   0
      Clean.Cup    0
```

- After find the missings values they can be replaced by find the mean, median if the columns are numerical in nature.
- And mode can be applied to the categorical data.
- Outliers can be visually seen by using the boxplot and scatter plot

```
[322]: sns.boxplot(x=df['Aroma'])  
[322]: <Axes: xlabel='Aroma'>
```



```
[323]: sns.boxplot(x=df['Flavor'])  
[323]: <Axes: xlabel='Flavor'>
```




Data filtering

Filtering data means choosing or not choosing certain information from a set of data using a set of criteria. This is important for finding important data, getting rid of unnecessary information, and improving the overall quality of the data.

`Dropna()`, `groupby()`, etc can be used to get particular columns required to be used for further analysis.

```
D1=data.groupby(['Variety']).size().reset_index(name='count').rename(columns={'Variety': 'Variety'})
```

	Variety	count
0	Arusha	6
1	Blue Mountain	2
2	Bourbon	226
3	Catimor	20
4	Catuai	74
5	Caturra	256
6	Ethiopian Heirlooms	1
7	Ethiopian Yirgacheffe	2
8	Gesha	12
9	Hawaiian Kona	44



The `quantile ()` method calculates the quantile of the values in a given axis. Default axis is row.

By specifying the column axis (`axis='columns'`), the `quantile ()` method calculates the quantile column-wise and returns the mean value for each row.

```
Q1=df.quantile(0.25)
Q3=df.quantile(0.75)
IQR=Q3-Q1
print(IQR)
```

```
Number.of.Bags      260.00
Aroma               0.33
Flavor              0.42
Aftertaste          0.41
Acidity             0.34
Body               0.34
Balance             0.42
Uniformity          0.00
Clean.Cup           0.00
Sweetness           0.00
Cupper.Points       0.42
Total.Cup.Points    2.50
Moisture            0.02
Category.One.Defects 0.00
Quakers             0.00
Category.Two.Defects 5.00
altitude_low_meters 450.00
altitude_high_meters 500.00
altitude_mean_meters 500.00
dtype: float64
```

```
file_cleaneddf=~((dfc ((Q1 - 1.5 * IQR)) | (df > (Q3 + 1.5 * IQR))).any(axis=1))
file_clean
```

	Number.of.Bags	Aroma	Flavor	Aftertaste	Acidity	Body	Balance	Uniformity	Clean.Cup	Sweetness	Cupper.Points	Total.Cup.Points	Moisture	CategoryOne
28	275.0	8.17	8.08	8.08	8.00	8.08	8.00	10.0	10.0	10.0	8.25	86.67	0.10	
44	12.0	8.08	8.08	8.00	8.00	7.83	8.08	10.0	10.0	10.0	8.00	86.08	0.11	
53	1.0	8.17	8.17	7.92	8.08	7.83	7.75	10.0	10.0	10.0	8.00	85.92	0.12	
57	150.0	7.83	8.00	8.00	8.17	7.83	8.00	10.0	10.0	10.0	8.00	85.83	0.12	
62	275.0	7.92	8.17	8.00	7.92	7.75	7.83	10.0	10.0	10.0	8.00	85.58	0.10	
1185	15.0	7.58	7.00	6.75	6.92	7.00	6.92	10.0	10.0	10.0	7.00	79.17	0.11	
1187	275.0	7.25	7.17	6.75	7.25	7.00	6.92	10.0	10.0	10.0	6.83	79.17	0.11	
1199	275.0	7.00	7.00	7.00	7.00	7.00	7.00	10.0	10.0	10.0	7.00	79.00	0.11	
1209	20.0	7.25	6.83	6.83	7.00	7.17	7.00	10.0	10.0	10.0	6.67	78.75	0.14	
1218	250.0	7.00	7.00	6.75	7.33	6.92	6.92	10.0	10.0	10.0	6.67	78.58	0.13	

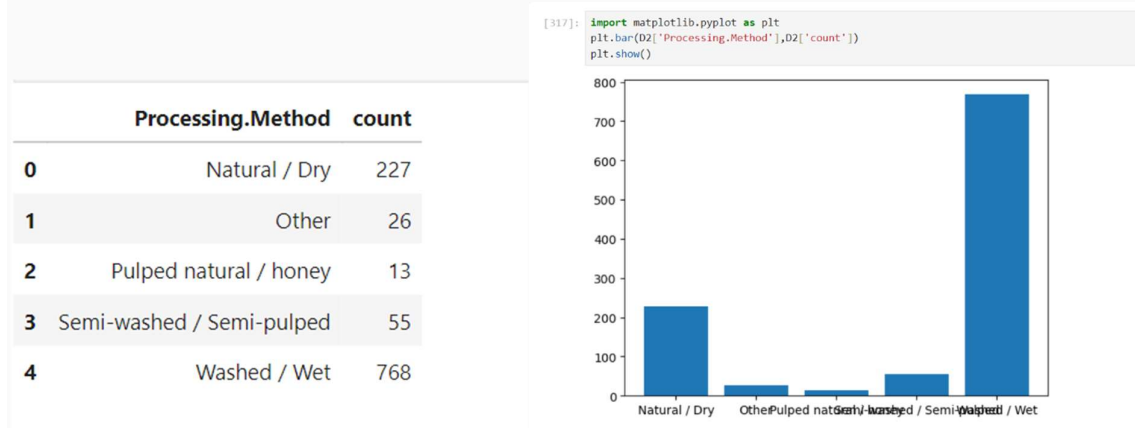
511 rows x 19 columns

After filtering and cleaning the data it is now ready to do further analysis needed .

Univariate analysis

Univariate data is a term used in statistics to describe data that consists of observations on only one characteristic or attribute. There is only one variable in

univariate data. Univariate data describes the variable's response pattern. For example, the analysis could look at a variable such as "age," "height," or "weight."



Bivariate analysis

Bivariate analysis is a statistical method to determine if there is a statistical link between the two variables and, if so, how strong and in which direction that link is. It is a helpful technique for determining how two variables are connected and finding trends and patterns in the data. In [statistical analysis](#), distinguishing between [categorical data and numerical data](#) is essential, as categorical data involves distinct categories or labels, while numerical data consists of measurable quantities.

Examples of bivariate analysis

Scatterplots

Correlation

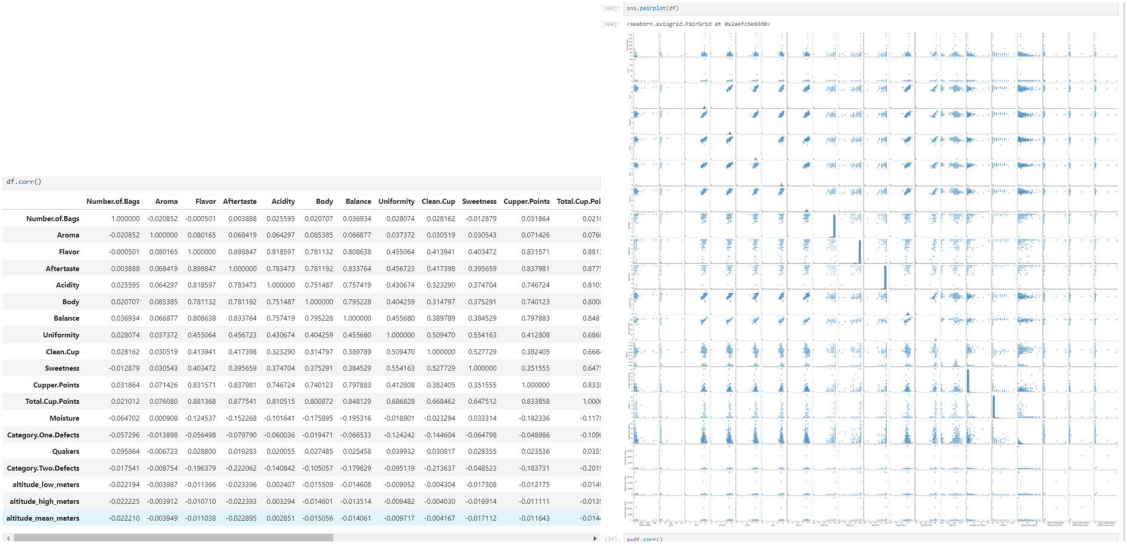
Regression

Chi square

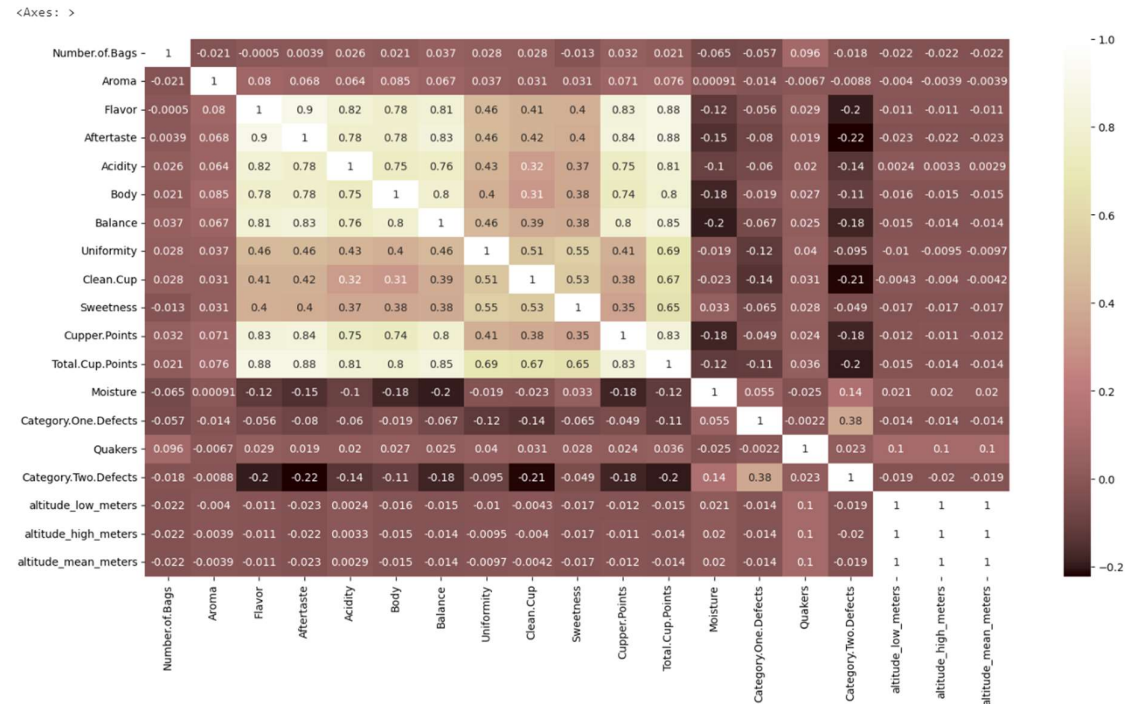
T test and ANOVA(Analysis of variance)

Correlation

Correlation is a statistical measure that shows how strong and in what direction two variables are linked. A positive correlation means that when one variable goes up, so does the other. A negative correlation shows that when one variable goes up, the other one goes down.



```
a=df.corr()
plt.figure(figsize=(18,9))
sns.heatmap(a,annot=True,cmap="pink")
```



Insight obtained from the Analysis

- Category 2 defects are more compared to category 1 defects.
- Washed/Wet processed beans has more defects.
- November, April & March Graded Beans has more defects.
- Beans with Moisture Range more than 10% has more defects.
- Harvest year 2021/2022 & 2022/2023 has more defected beans.

Note : Here, the Coffee Bean production in the respective years and Processed method is high, so it may also be the reason for more def